

Computer Programming – CS 6011

Lecture 16: Threads

MASTER OF SOFTWARE DEVELOPMENT (MSD) PROGRAM

J. DAVISON DE ST. GERMAIN

FALL 2022

Miscellaneous

- ▶ Questions?
 - ▶ Lambda Functions
 - ▶ Anonymous / unnamed / in-line functions
- `() -> { }`
- ▶ HTML Injection
 - ▶ ``

Lecture 16 – Topics

3

- ▶ Threads
 - ▶ Executing multiple pieces of code at the same time.
 - ▶ Synchronization

Processes vs Threads

- ▶ A process is a running program
 - ▶ `% ps -def | more`
- ▶ A single program, in order to do multiple things at the same time, can use multiple threads.
- ▶ Processes cannot “talk” to each other (directly*).
 - ▶ Threads can (as we will see).
- ▶ Definition: A thread, or thread of execution, is a mechanism that allows a program (a process) to divide itself into two or more simultaneously running tasks. (~Wikipedia)

Why Threads

5

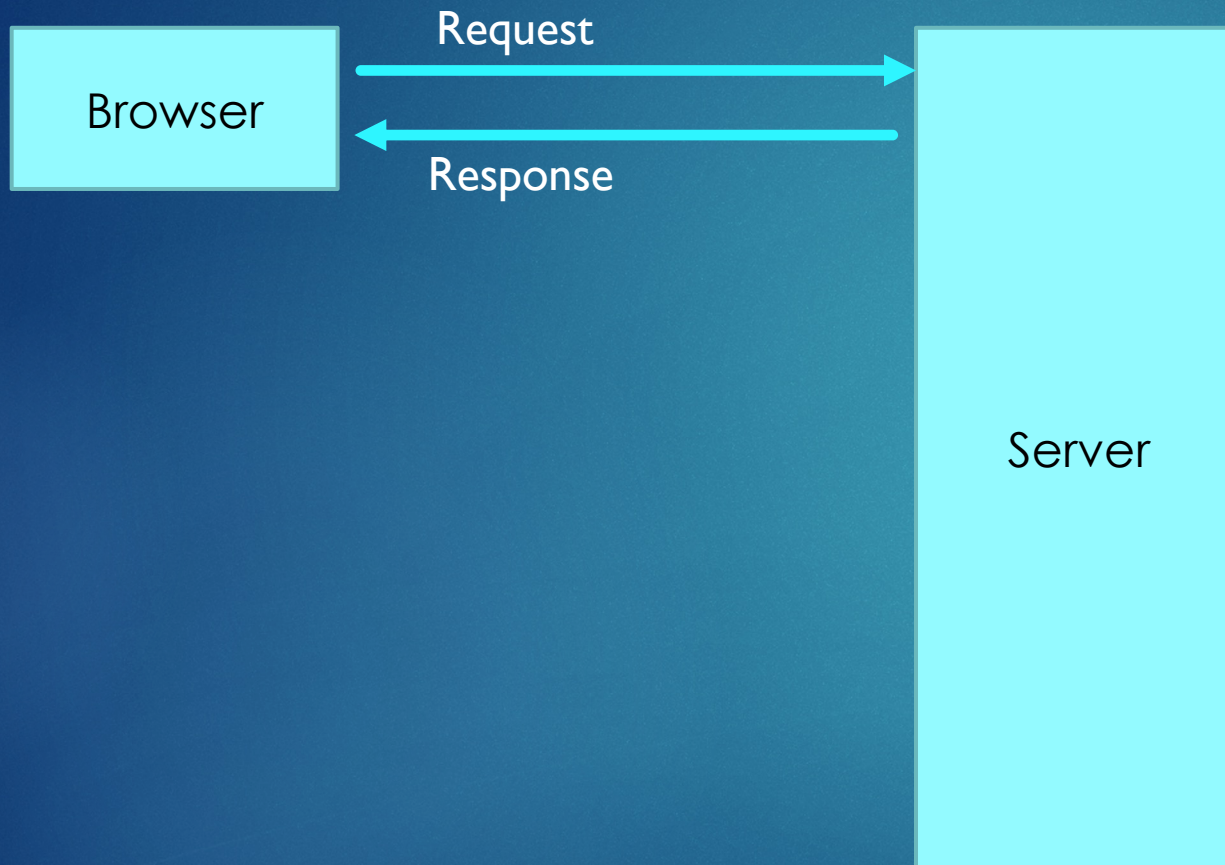
- ▶ (Many) Programs need to do multiple things concurrently.
- ▶ CPUs stopped getting faster* a while back... so instead of a faster CPU, the hardware people give us more CPUs (cores)...
 - ▶ Threads are one way (of several) a computer can make use of multiple cores.
 - ▶ If this laptop has 4 cores – how could I use them all at the same time?
 - ▶ I could run 4 different programs simultaneously.
 - ▶ Threads allow a single program to use multiple computer resources (cores) at the same time.
- ▶ **Concurrency** – running multiple threads at the same time. Also called **Parallelism***.
 - ▶ We see this all the time – most explicitly in User Interfaces.
 - ▶ While a program is doing something, you can still interact with the UI.
- ▶ Any questions on this concept? Where have we seen these things so far in our projects?

Thread we have “Seen” before.

- ▶ The playing of a clip by the synthesizer... how did we “wait” until the sound was done before doing anything.
 - ▶ `while (!done) { }`
 - ▶ What happened to our program during this time?
 - ▶ It became frozen.
 - ▶ Added a listener...
 - ▶ The audio library actually creates its own thread and notifies the main program (thread) when it is done.
- ▶ As we work with threads, we will see why it is important to use them, and to use them correctly. For example:
 - ▶ In some of your game projects from CS 6010 you added while loops to check for events. These loops did not stop until the user pressed a button or clicked the mouse.
 - ▶ The program became unresponsive (and the spinning beachball of death appeared).

Web Server

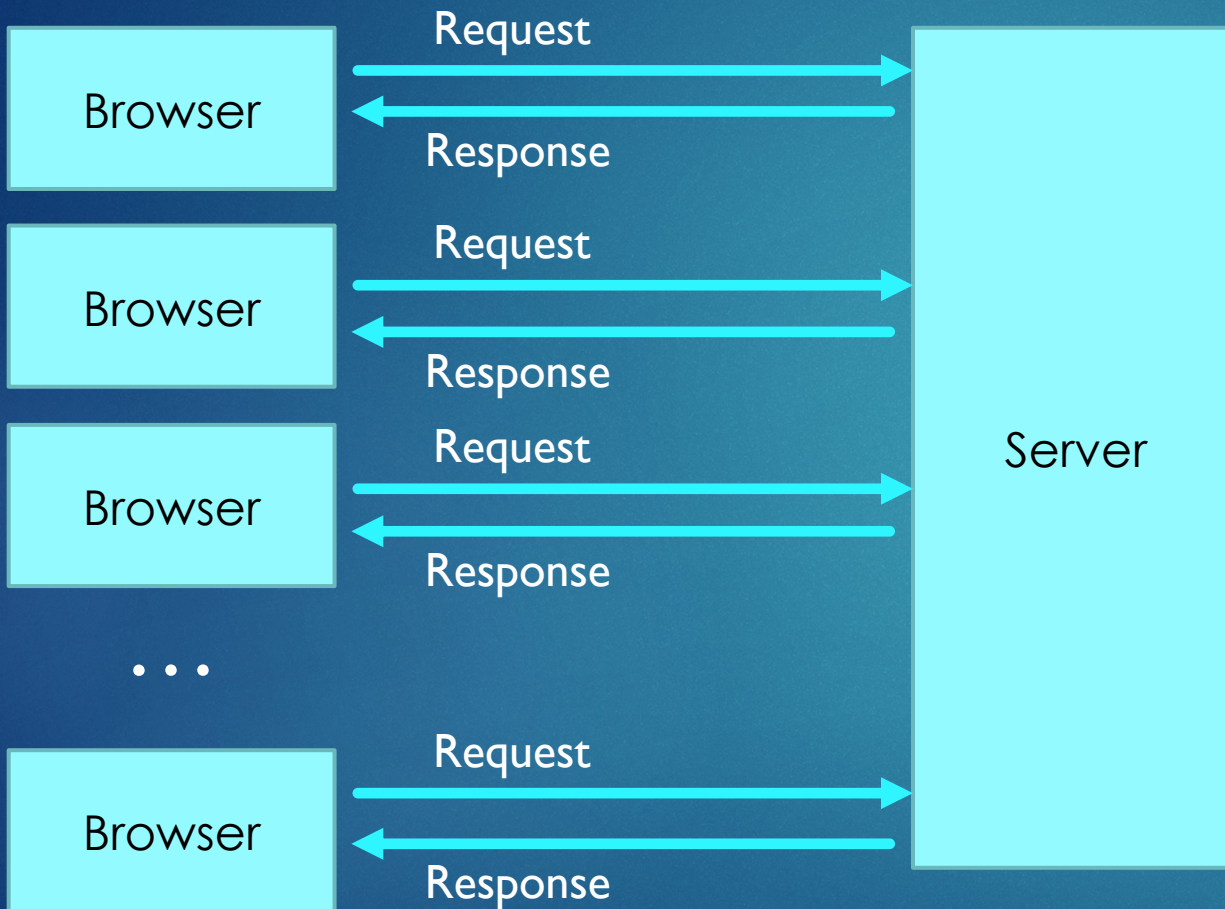
7



- ▶ With a single client (browser), the
 - ▶ Request
 - ▶ Process Request
 - ▶ Response
- ▶ Happens in an apparent *synchronous* manner.
- ▶ But what would happen if there were a lot of clients making requests at the same time?

Parallel Web Server

8



- ▶ Multiple requests hit server simultaneously...
 - ▶ Is this possible?
 - ▶ Technically no – network card only handles one message at a time (though very fast).
- ▶ Multiple requests can hit server very quickly (before it can completely handle them).
- ▶ What does the server do?
 - ▶ Could handle them in the order they arrived, one at a time (eg: checkout line at the grocery store)... or
 - ▶ Create a thread to handle each one.
 - ▶ Threads (can) run on their own hardware (Core), and thus can respond to the requests simultaneously. Well, at least process the request simultaneously.

What is a Thread?

- ▶ Basically a running function – one that runs independently (and usually at the same time) as other code.
 - ▶ You can have many of these Threads (methods) running at the same time. Each thread doing something for you.
- ▶ This allows for (one type of) *asynchronous* programming.
- ▶ What resources does a computer need to run a function?
 - ▶ Memory – Call Stack
 - ▶ CPU Time – Either on the same core or on a separate core from other Threads / Processes.
- ▶ Threads run in the same “memory space” (heap) as each other.
 - ▶ This means that if you don’t play attention, one thread can alter the same variables that another thread is using.

Threads in Java

10

- ▶ It is easy to make threads in Java – in fact your `main()` function is actually running in a thread.
- ▶ Threads are created using the `Thread` class.
- ▶ However, Threads use a “Runnable” object to do the actual work.
 - ▶ Runnable is an interface that contains one function: `void run()`
 - ▶ You *implement* a Runnable and provide the `run()` function.
 - ▶ Then you pass the runnable to the thread (via its constructor) for it to use.
- ▶ Important Thread methods:
 - ▶ `start();` // Start running the thread (in parallel with any other threads).
 - ▶ `join();` // Wait for a thread to finish running.
 - ▶ `threadId();` ~~`getId();`~~ // Returns a unique identifier for the Thread object.
 - ▶ `Thread.currentThread();` // Returns the Thread object the current method is running in.

Creating a Thread (Lambda Version)

11

- ▶ Creating a Thread as a *lambda* function. (le, in-line)

```
Thread t = new Thread( () -> {  
    // Code to be run in this thread.  
    // Many times the code is within a while() loop so that it continues  
    // doing work for us..  
} );
```

t created

- ▶ To actually get the thread running, we need to use the `start()` method.

```
t.start();
```



Ending a Thread...

- ▶ If `main()` can't do anything else until `t` finishes its jobs, we use:

```
t.join() // Main waits for t to finish...
```

- ▶ Once `t` finishes, it joins back with `main` and we are back to a single thread, and `main()` continues to execute.
- ▶ `main` is *blocked* (does not execute any code) after calling `t.join()`.



(Mis-)Synchronization

13

- ▶ Think about two people working together – a customer and a baker.
 - ▶ Each of these “people” is part of our program and represented by a different thread (so that they can be doing their own things at the same time).
 - ▶ The customer asks for some cookies... [I would like one hundred cookies]

Customer: I would like one hundred cookies.

Baker: Okay, baking one cookie.

How do Threads Talk to Each Other?

- ▶ Messages – We'll talk about this in a future class.
- ▶ Shared Memory
 - ▶ Common variables
 - ▶ One Whiteboard
 - ▶ Everyone is a ghost
 - ▶ No one knows when anyone else is doing something
 - ▶ Everyone has a marker
 - ▶ What can go wrong with this?
 - ▶ Two threads read / write at the same time.
 - ▶ Example, two threads using / updating a distance
 - ▶ Read currentDistance
 - ▶ Calculate the new value for currentDistance
 - ▶ Store new value for currentDistance
- ▶ Another Thread
 - ▶ Changes value of currentDistance
- ▶ *Critical Section*
 - ▶ Can't be interrupted by another thread or problems occur.
 - ▶ Mutexes / Locks are used to control access to critical section.
 - ▶ Java uses *synchronized* (keyword)

Critical Section Example (Not Protected)

15

▶ `int x = 33;`

▶ If we aren't careful, this happens:

▶ Thread 1:

```
int i = x + x;
```

▶ What should `i`'s value be equal?

▶ How was that calculated? What's the value of `x`?

▶ Let's do this one step at a time...

▶ The computer (well, an individual thread) can only do one thing at a time, so what is the first thing it does in the line above? What happens first?

▶ Computer loads the value of `x` from memory. Now what?

▶ Computer is about to load the value of the 2nd variable (which just happens also to be `x`).

▶ But wait... there is a 2nd thread that is running:

▶ Thread 2:

```
x = 0;
```

▶ Now back to Thread 1...

▶ `i` is set to 33...

Synchronization

16

- ▶ Because threads are running at the same time (and in a nondeterministic order), any data (variable, etc) that is read (accessed) by one thread, but was written (assigned) by another thread is a potential problem.
 - ▶ These types of bugs happen “randomly” and are very hard to reproduce.
 - ▶ The best approach to fixing them is to prevent them in the first place.
- ▶ Synchronization is an approach to prevent different threads from accessing the same data at the same time.
- ▶ We'll learn several ways to safely share data between threads.

Java Synchronized

17

- ▶ Synchronized Method

```
public class MyClass {  
    public synchronized void doit() { ... }  
}
```

- ▶ Each thread that calls `doit()` – *on the same object* – checks to see if any other thread is currently running the `doit()` code, and if so, waits until the other thread is done.

```
MyClass myObj = new MyClass();
```

- ▶ What if two different threads have access to `myObj` and try to run `doit()`?

- ▶ Thread t1

```
myObj.doit();
```

- ▶ Thread t2

```
myObj.doit();
```

- ▶ Order of t1 `doit()` vs t2 `doit()`?

- ▶ Undefined... but only **one** will execute at a time.

Java Synchronized

18

- ▶ Synchronized Method

```
public class MyClass {  
    public synchronized void doit() { ... }  
}
```

- ▶ However, threads using *different* objects can access the same code simultaneously!

```
MyClass myObj1 = new MyClass();
```

```
MyClass myObj2 = new MyClass();
```

- ▶ Thread t1

```
myObj1.doit();
```

- ▶ Thread t2

```
myObj2.doit();
```

- ▶ Order of t1 doit() vs t2 doit()?

- ▶ Undefined... but both can execute at the same time – they are using *different* objects!

General Rule

- ▶ When we have shared objects (such as `myObj` on slide 17), we can use synchronized methods to allow multiple threads to call the same function (at the “same” time) on that object.
 - ▶ This will protect shared variables (data) from being corrupted.
 - ▶ But will only allow one thread at a time to actually run...
 - ▶ In other words, the program has been *serialized* during this part of its execution and is thus not enjoying any speed up benefits.

Threading is Easy... If:

20

1. None of the threads share any data; or
2. All the shared data is *read-only*.

Threading Example

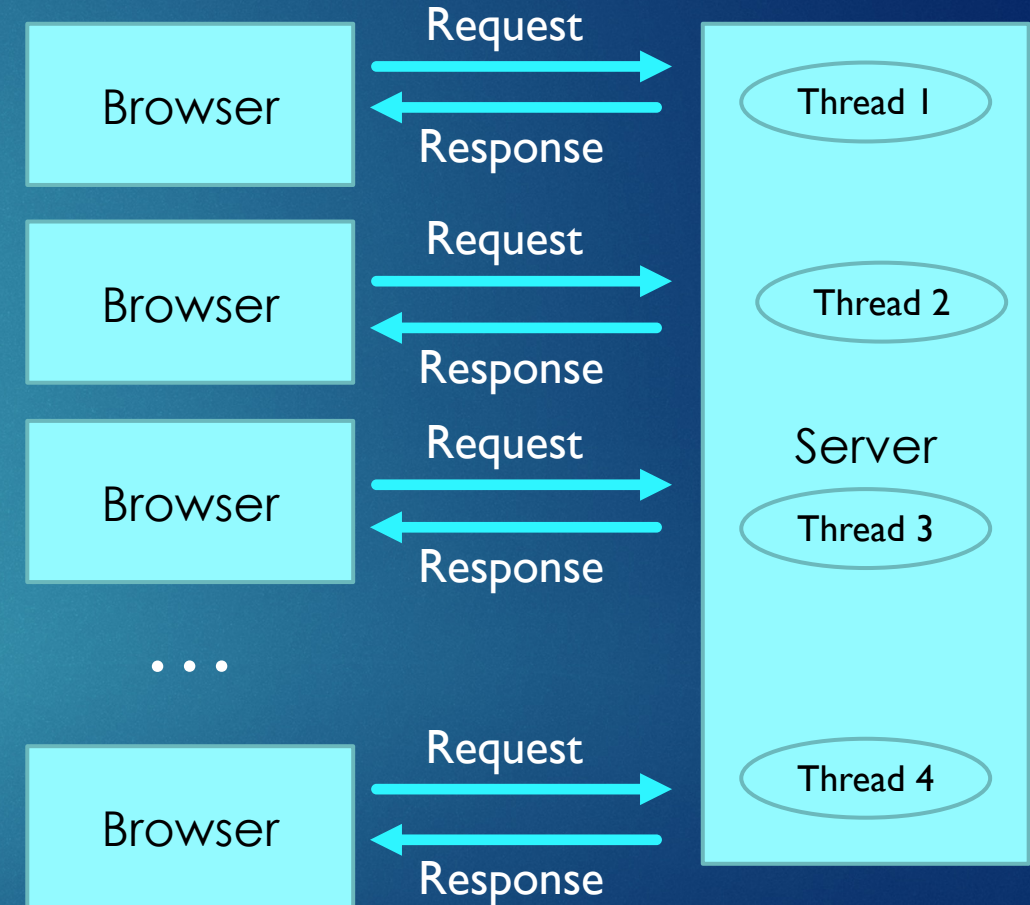
- ▶ Create a *Runnable* class.
 - ▶ Runnables only have one method: *run()*.
- ▶ Create threads to execute the code (in parallel).
 - ▶ `start()` the thread and off it goes...
 - ▶ Executing the `run()` method.
- ▶ In-class example...

- ▶ Examining a “critical region”.
- ▶ Basically $x = x + \text{number}$;
- ▶ What happens when a bunch of threads try to help each other to do this (quickly) at the same time?

Assignment – Multithreaded Web Server

23

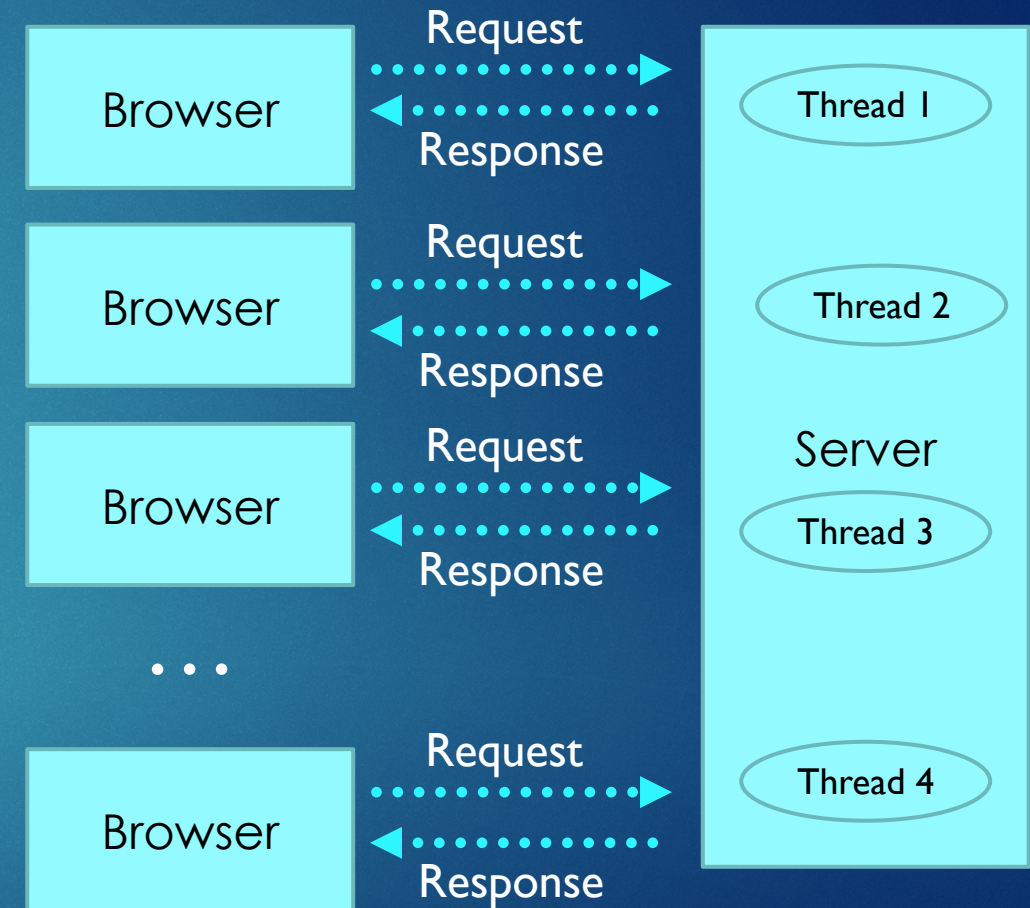
- ▶ Why would we want to do this?
 - ▶ So multiple clients can talk to the server at the same time.
- ▶ When do we want to create a thread?
 - ▶ When a new client has connected to us.
 - ▶ When does this occur?
 - ▶ `ss.accept()`
- ▶ When do we destroy a thread?
 - ▶ When the server is done responding to a client.
 - ▶ Do we have to do this ourselves?
 - ▶ No – the garbage collector will take care of them for us.
- ▶ Recap – What happens now...



Assignment – Multi-Threaded Web Server

24

- ▶ At this point Thread 1 – 4 have finished running (ie: finished processing and responding to the requests)... The `run()` function ends and thus nothing is referencing them anymore... so:
 - ▶ The garbage collector kicks in and...
 - ▶ The Thread objects are gone.
- ▶ Considerations for threading...
- ▶ Do you have any shared data for your webserver?
 - ▶ Do you have static variables?
 - ▶ As a side note, we do have a “shared” resource. What is it?
 - ▶ Files
 - ▶ But they are “read-only”* so don’t need to worry about this.



Multi-Threaded Web Server Demo

- ▶ Loading multiple pages / images at the same time.
 - ▶ I'm using an (artificial) random delay before serving each image.
- ▶ Network tab on browser tools...

Tuesday Assignments

- ▶ Code Review – Web Chat Client
- ▶ Lab – Threads
- ▶ Assignment – Multi-Threaded Web Server

~ Fin ~