

Computer Programming – CS 6011

Lecture 17: Web Socket Implementation

MASTER OF SOFTWARE DEVELOPMENT (MSD) PROGRAM

J. DAVISON DE ST. GERMAIN

FALL 2023

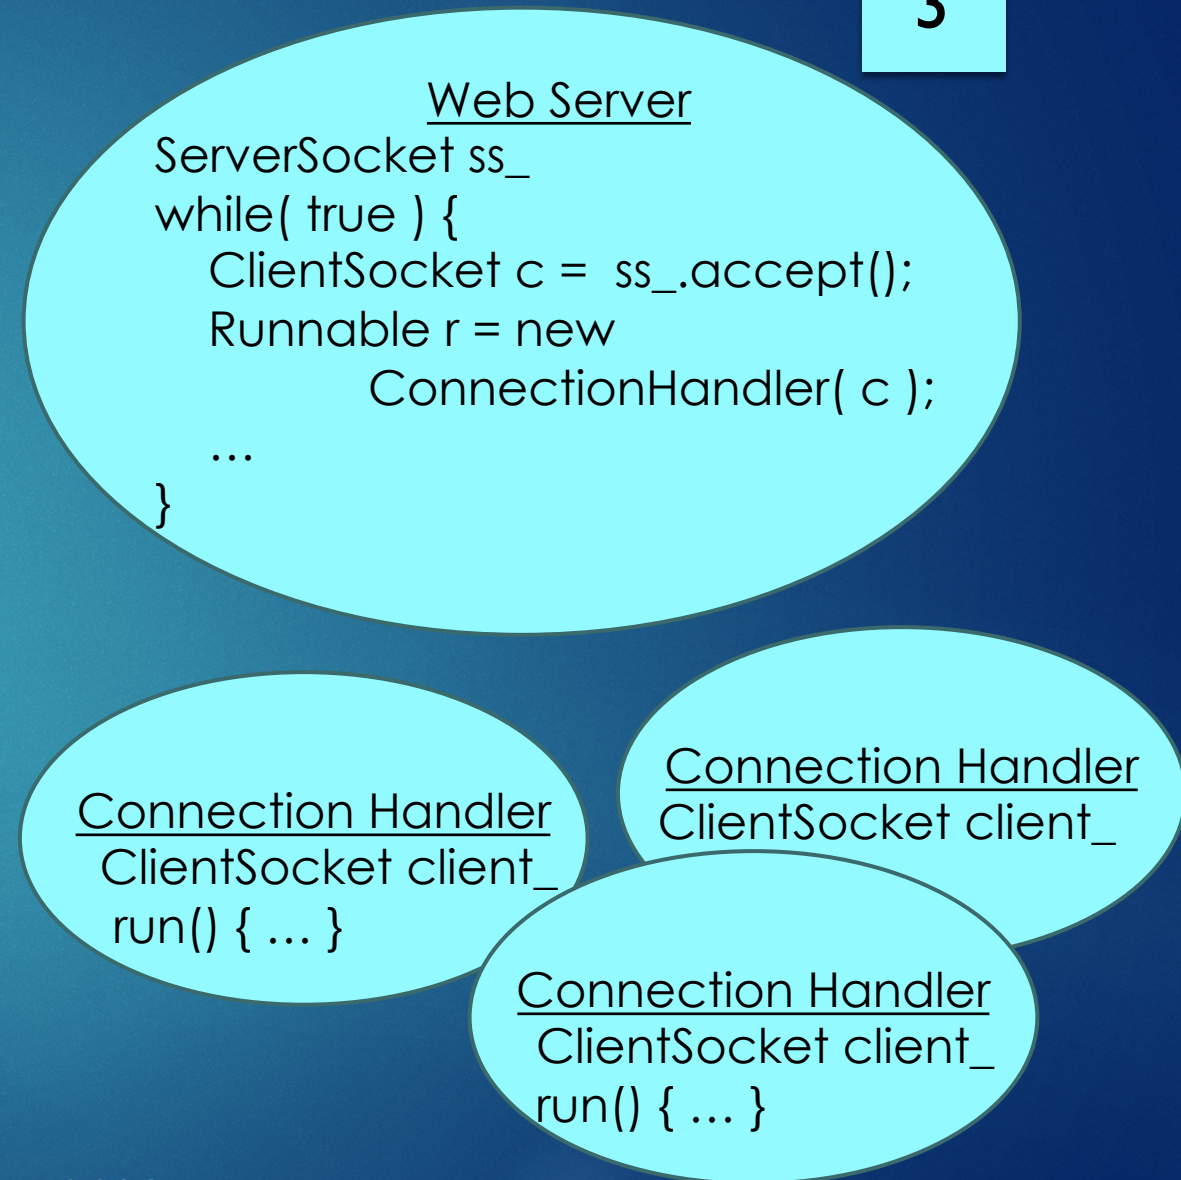
Miscellaneous

- ▶ Passing Data to a Thread (Runnable)
 - ▶ Constructor
 - ▶ Global Variables
 - ▶ Starting / Joining
- ▶ Questions?

Web Server Classes

3

- ▶ How to organize the data / functionality of your web server? (Bubble Diagram!)
 - ▶ What are the main pieces?
 - ▶ The server itself.
 - ▶ A connection handler...
 - ▶ The ability to handle (parse) a request
 - ▶ The ability to send a response.
 - ▶ Given those pieces, what data does the server need?
 - ▶ A ServerSocket.
 - ▶ What data does a “connection handler” need?
 - ▶ The ability to read / write data to the client... so what data?
 - ▶ The Client Socket
 - ▶ Which of these objects should be the thread (or more specifically, the runnable)?
 - ▶ Connection Handler



Sharing Data Among Objects...

- ▶ Who creates (gets) the Client Socket?
 - ▶ The Server
- ▶ Who needs the client socket in order to do its job?
 - ▶ The Connection Handler (Runnable).
- ▶ Who creates a Connection Handler?
 - ▶ The Server
- ▶ How does the Connection Handler get the client socket?
 - ▶ The server gives the client socket to it... how?

```
class Server { // Pseudo Code
```

```
private ServerSocket ss = ...;
```

```
public void main() {
```

// How does the server get the client socket?

```
Socket client = ss.accept();
```

// How does it give it to the connection Handler?

```
ConnectionHandler ch = new ConnectionHandler( client );
```

```
// Use the constructor -----^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^.
```

// How to make this ConnectionHandler go do its thing?

```
Thread t = new Thread( ch )
```

```
t.start();
```

```
// Shove it into a thread and start it running.
```

- ▶ Notice, no `HttpRequest` or `HttpResponse` objects in the above code... Why not?
 - ▶ The `ConnectionHandler` uses them, not the `Server`!

Connection Handler // Pseudo code

5

```
class ConnectionHandler implements Runnable {  
    // In order to read headers / send a response, what does the CH need?  
    Socket client_;  
    // How does the CH get the client socket?  
    public ConnectionHandler( Socket client ) {  
        // what code goes here?  
        client_ = client;  
    }  
    // Methods?  
    void run() {  
        // We use client_ for the following:  
        // what happens here (high level)?  
        // read / parse request headers  
        // send response headers / file to requestor  
    }  
}
```

// How do we read / parse headers in one line of code?
request = new HttpRequest(client_);
request.parse(); *// ok, maybe 2 lines...*

// How to send a response (again, with only one line)?
response = new HttpResponse(client_, request.getFilename() , request.getHeaders())



Lecture 17 – Topics

6

- ▶ Web Sockets
 - ▶ Implementing a WebSocket (Chat) Server
 - ▶ We'll be discussing this over the next two days...

WebSocket Protocol

7

- ▶ First, we need to understand the protocol.
 - ▶ HTTP Request followed by a HTTP Response...
 - ▶ This is called a *handshake*.
 - ▶ What header fields are used, and how are they set?
 - ▶ WebSockets then switch to using binary data packets.
 - ▶ What does this look like?
- ▶ Look at these documents... can you answer the above questions?
 - ▶ https://developer.mozilla.org/en-US/docs/Web/API/WebSockets_API/Writing_WebSocket_servers
 - ▶ <https://datatracker.ietf.org/doc/html/rfc6455>

Some Key Points From The Docs...

8

- ▶ From: <https://datatracker.ietf.org/doc/html/rfc6455>
- ▶ 1.1 Background
 - ▶ Historically, creating an instant messenger chat client as a Web application has required an abuse of HTTP to poll the server for updates while sending upstream notifications as distinct HTTP calls.
- ▶ 1.3 Handshake
 - ▶ Important header fields?
 - ▶ Request
 - ▶ Sec-WebSocket-Key: dGhlIHNhbXBsZSBub25jZQ==
 - ▶ Response:
 - ▶ Sec-WebSocket-Accept: s3pPLMBiTxaQ9kYGzzhZRbK+xOo=
- ▶ 5.2 Base Framing Protocol
 - ▶ The specification of a WebSocket message (Header + Payload)

Starts with HTTP Protocol

9

- ▶ What is the first thing that happens when a client tries to establish a WebSocket connection to a server?
 - ▶ A Handshake occurs... in other words, the client:
 - ▶ Sends an HTTP Request... So we need to parse the request for the relevant information:
 - ▶ How do we (or did you) store the request fields?
 - ▶ Map
 - ▶ HashMap (the actual implementation of a Map) – using what type of Keys / Values
 - ▶ HashMap< String, String >
 - ▶ Which key(s) do we look at to decide what type of request this is?
 - ▶ Connection: Upgrade
 - ▶ Sec-WebSocket-Key: dGhlIHhnbXBsZSBub25jZQ==

```
GET /chat HTTP/1.1
Host: example.com:8000
Upgrade: websocket
Connection: Upgrade
Sec-WebSocket-Key: dG...ZQ==
Sec-WebSocket-Version: 13
```


Responding to a Request

- ▶ Remember, HTTP Requests coming to your server can be normal requests (like what we have handled already), or a WebSocket request.
 - ▶ Once we have determined which type of request it is (previous slide) – if it is a WebSocket...
- ▶ Handshake Request we will need to send back the required response (as a normal HTTP Response)... [See Handshake on next slide]
 - ▶ and then...
- ▶ Switch to the WebSocket protocol for all future messages.
 - ▶ From plain text (ASCII) Request / Response, switch to a binary protocol.

```
void run() {  
    // We use client_ for the following:  
    // what happens here (high level)?  
    // read / parse request headers  
    // send response headers / file to requestor  
    // After handing a normal http request,  
    // what is the last thing we actually need to  
    // do here?  
    client_.close();  
}
```


WebSocket Handshake

- ▶ What response header (key) do we use to respond to the WS request?
 - ▶ `Sec-WebSocket-Accept: <value>`
- ▶ What `<value>` will it contain (ie, what value do we put in this field)?
 - ▶ Concatenate the client's `Sec-WebSocketKey` and the magic string:
 - ▶ `"258EAF5E-E914-47DA-95CA-C5AB0DC85B11"`
 - ▶ Then use:
 - ▶ SHA-1 Hash and Base64 encoding
 - ▶ `MessageDigest` class will do the SHA-1 hashing
 - ▶ `Base64` class will encode the output of the `MessageDigest` into a string.
- ▶ How to test that the Handshake worked?
 - ▶ Wireshark can show that the server is sending packets back.
 - ▶ But much easier... how does the browser know when a WebSocket handshake has completed?
 - ▶ It runs the `ws.onopen` callback.

```
GET /chat HTTP/1.1
Host: example.com:8000
Upgrade: websocket
Connection: Upgrade
Sec-WebSocket-Key: dG...ZQ==
Sec-WebSocket-Version: 13
```


↑ Handshake ↓

```
HTTP/1.1 101 Switching Protocols
Upgrade: websocket
Connection: Upgrade
Sec-WebSocket-Accept: s3p...o=
```


WebSocket Messages

12

- ▶ Server doesn't close the client socket, it continues to use it to send data...
- ▶ However, the messages now look like: (because this is a binary format)



- ▶ On client (web browser / JavaScript), it's easy:
 - ▶ `ws.send("my message");` // Client library formats message for us.
- ▶ On the server, it's (a little?) more complicated...
 - ▶ First, why do we need the `Header` for this message (mostly)?
 - ▶ It has the length of the payload. (This is common to almost all binary formats.)
 - ▶ Also a few other fields of interest...

Example WebSocket Message

13

- ▶ Data going across the network (representing a WebSocket message) looks like:

```
<- 0x82 0xFE 0x0100 0xABCD ...
```

- ▶ In Binary this looks like:

```
<- 1000 0010   1111 1110   0000 0001 0000 0000   1010 1011 1100 1101   .....
```

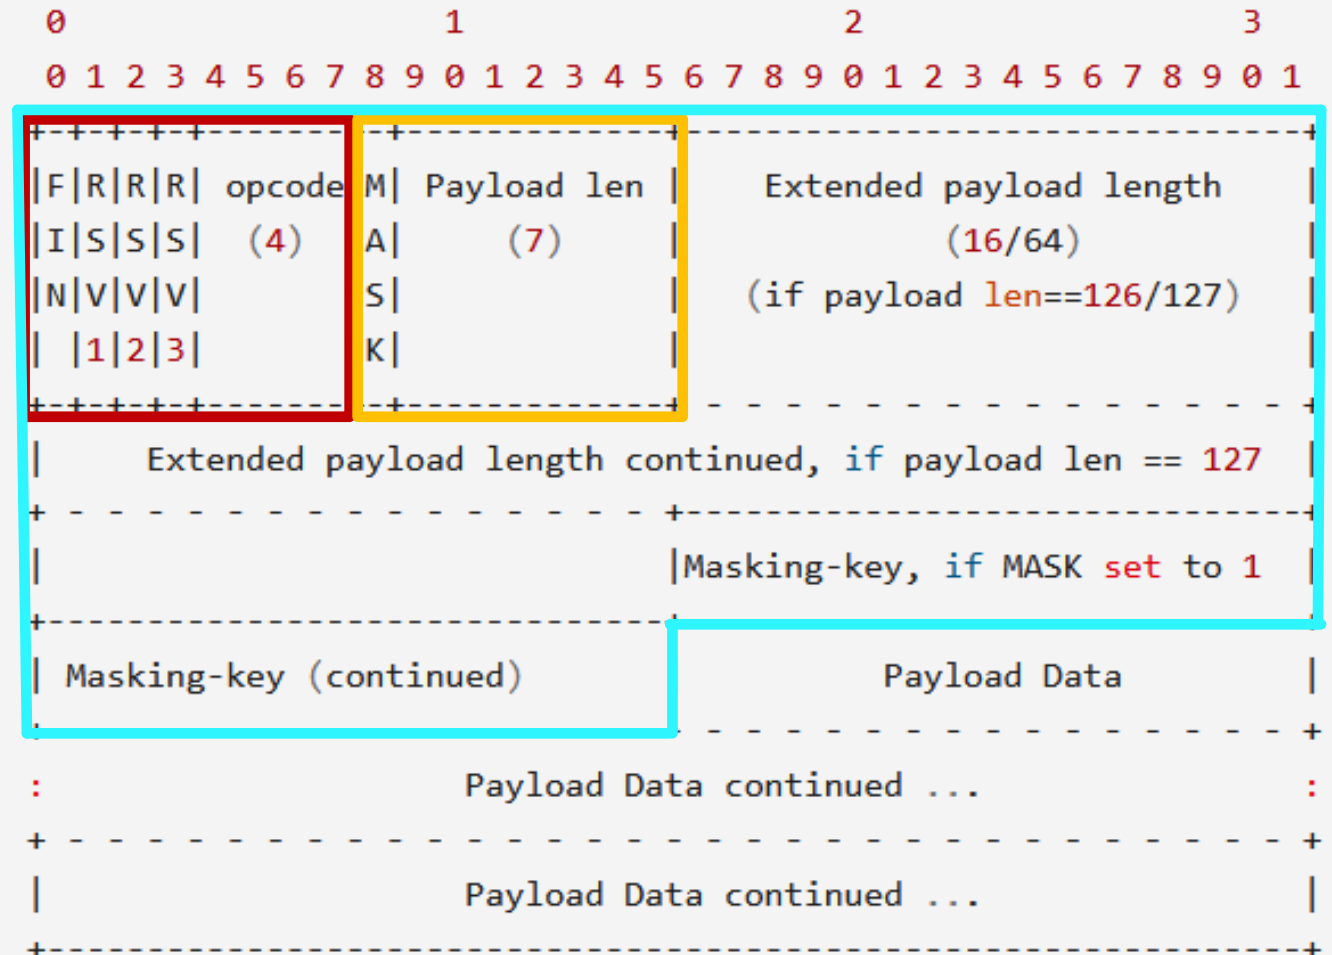
- ▶ What do each of those bits mean?
 - ▶ Let's figure that out starting on the next slide...

WebSocket Header

14

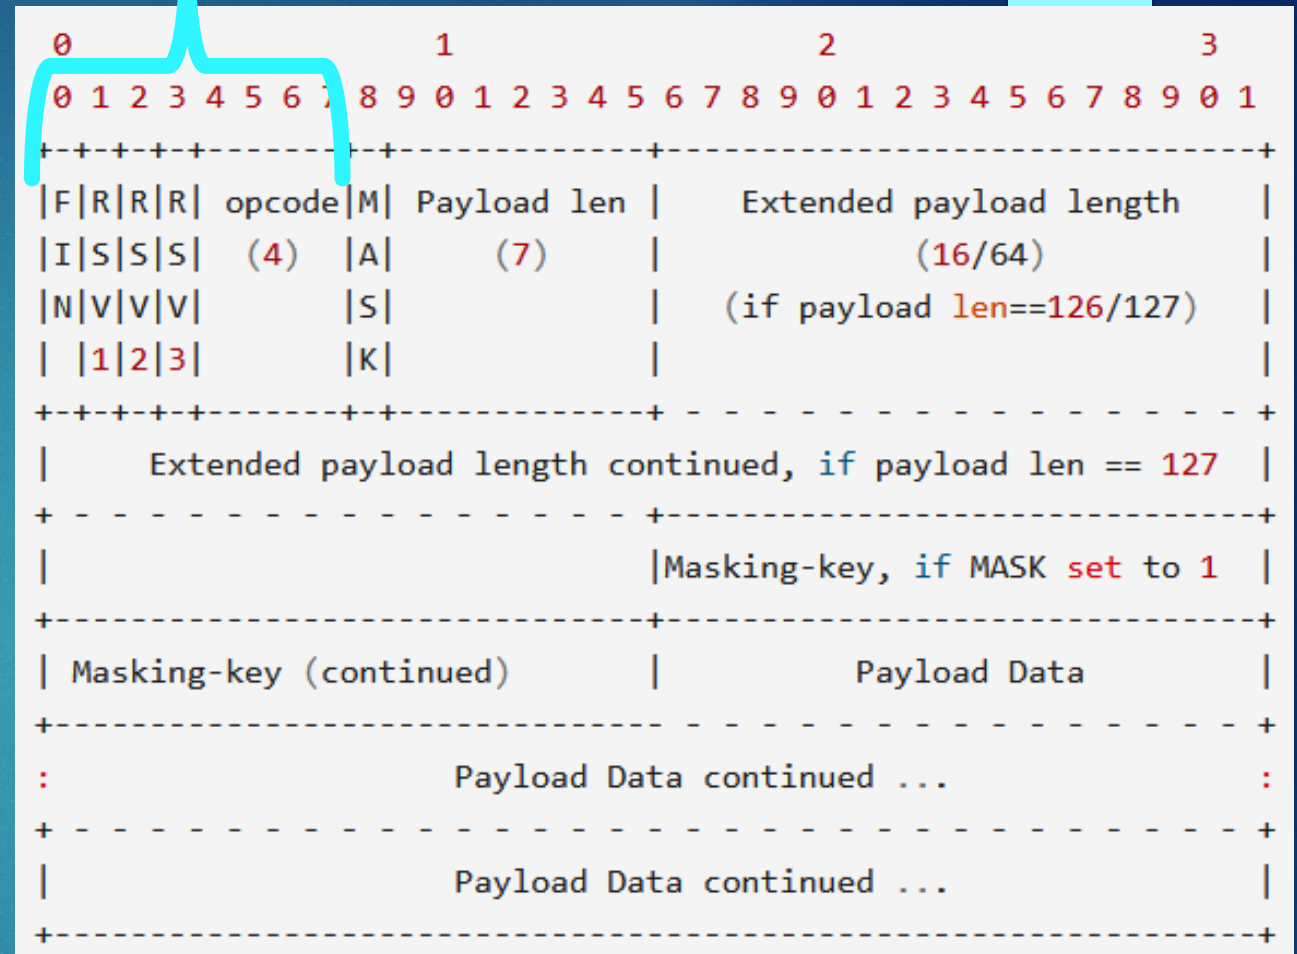
- ▶ Bytes are shown left to right.
- ▶ How big is the header?
 - ▶ Often 2 bytes, but...
 - ▶ Length data can be in **byte 1** (2nd byte), or be the next 2 (or 8) bytes
- ▶ Byte 0 (Shown In Red)
 - ▶ Opcode?
 - ▶ `b0 & 0x0F;`
- ▶ Byte 1 (Shown in Orange)
 - ▶ Masked?
 - ▶ `(b1 & 0x80) != 0;`
 - ▶ Also contains the length*.

Frame format:



WebSocket Header

- ▶ Max size of header?
 - ▶ 14 bytes
- ▶ Header Size:
 - ▶ If Payload is 0-125 bytes, with a client to server message:
 - ▶ Header size: 6 bytes (2 + mask)
- ▶ ...



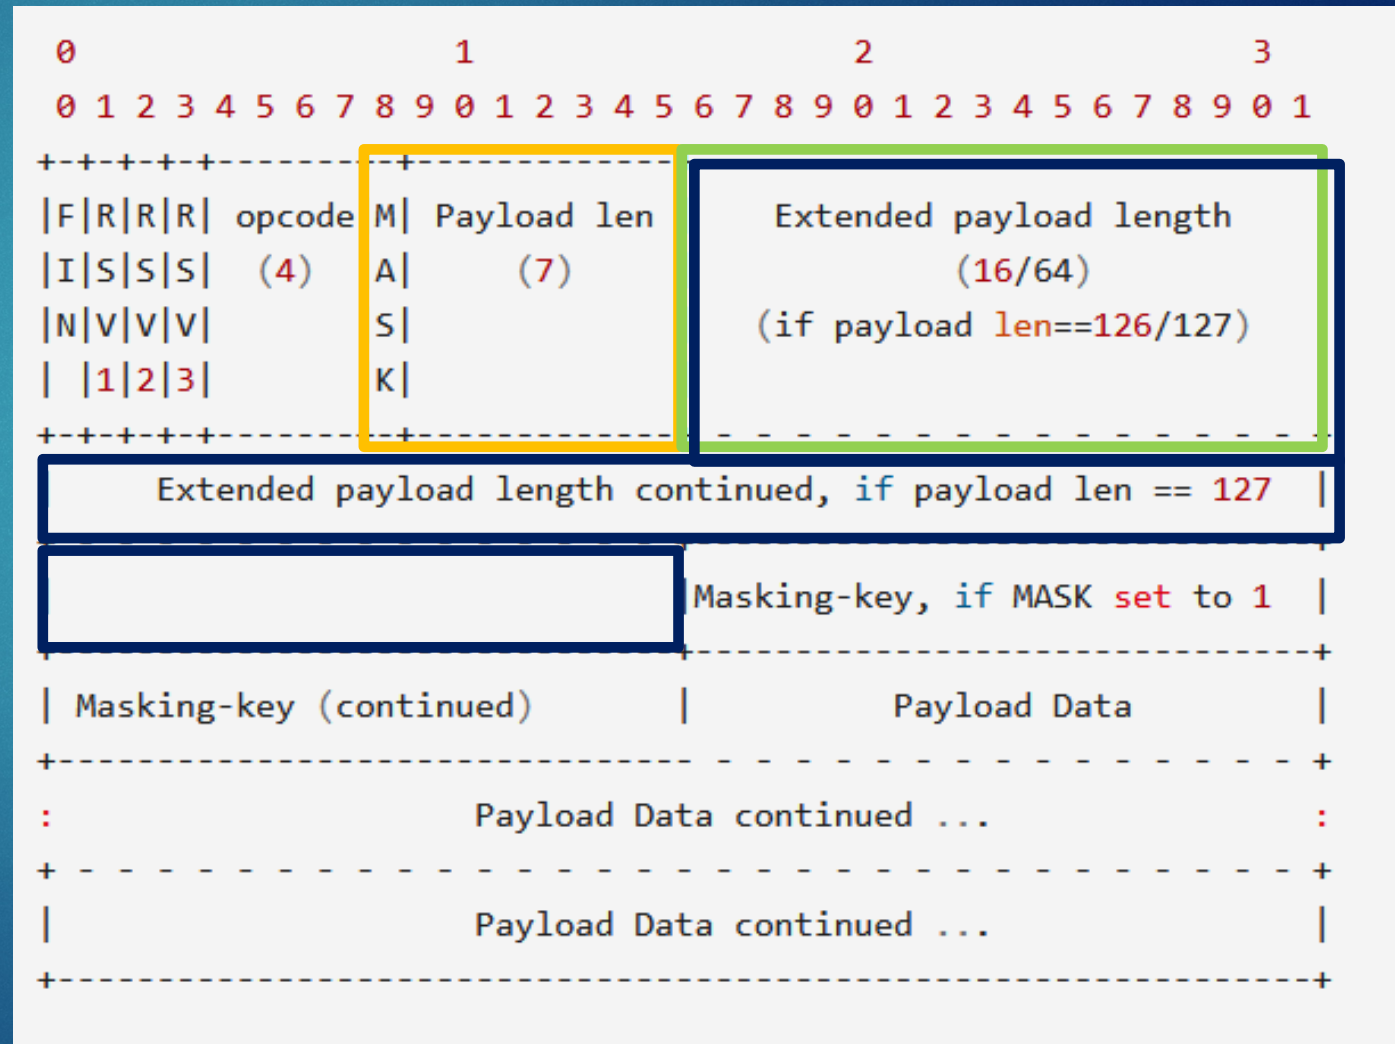
0	1	2	3	4	5	6	7	8	9	10	11	12	13
f op	m len	[opt	len]	[more	opt	ional	len	bytes	...]	[opt	ional	...	mask]	[Payl	oad..

First byte: Fin / Rsv / Rsv / Rsv / Opcode

WebSocket Header

16

- ▶ Header size changes based on?
 - ▶ Length of payload
- ▶ If payload length is 125 bytes or less, then the length is found in **b1** (remember to disregard the mask bit).
 - ▶ `lenGuess = b1 & 0x7F`
- ▶ If the value in **b1** is 126, then the length of the payload is found in bytes **b2-b3**.
 - ▶ Type? Range?
 - ▶ *Unsigned Short... why?*
- ▶ If the value in **b1** is 127, then the length of the payload is found in bytes **b2-b9**.
 - ▶ Type?
 - ▶ Long (Will this ever happen?)



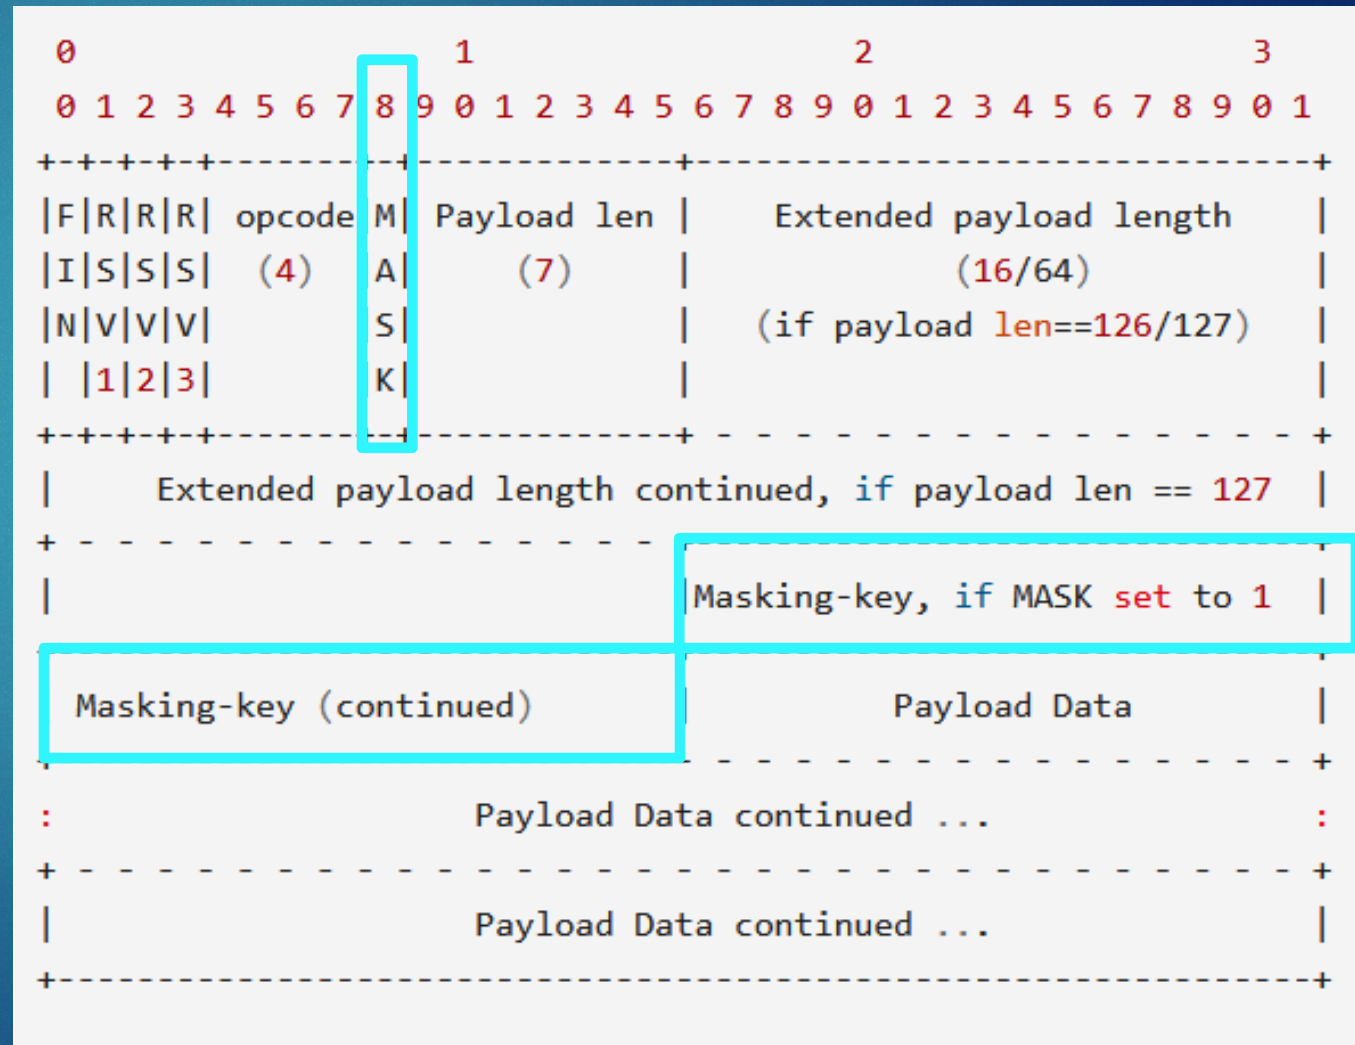
WebSocket Header

17

- ▶ Mask is only used for messages that go from? [RTFM*]
 - ▶ *Read the manual. 😊
 - ▶ Client to server.
- ▶ For these messages, the mask bit will be set to 1, and the masking-key will be 4 bytes.
- ▶ Conversely, messages sent from the server to the client are not masked, and there is NO masking-key.
- ▶ Code for unmasking is straightforward (and right out of the manual)...

```
var DECODED = "";
```

```
for (var i = 0; i < ENCODED.length; i++) {  
    DECODED[i] = ENCODED[i] ^ MASK[i % 4];  
}
```



Other Header Fields

18

▶ FIN bit?

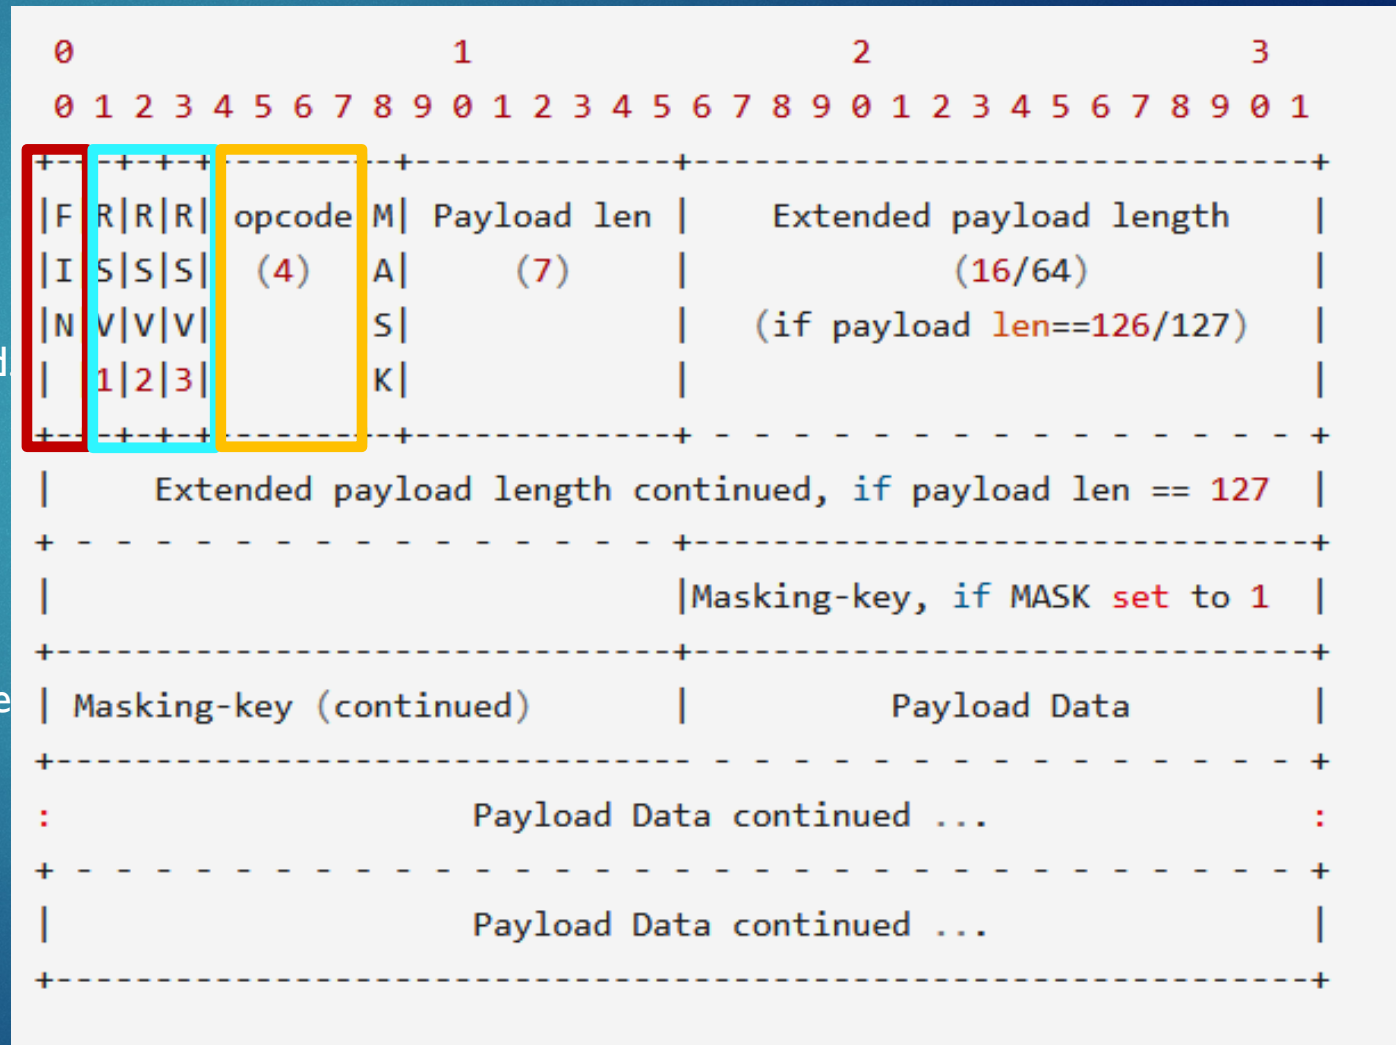
- ▶ WebSocket messages can be sent in parts, though we probably won't see these...
- ▶ FIN of 1 means this “series of messages” is over... Since we will only have one message in each series, this should always be 1.

▶ RSV – Reserved

- ▶ Bits reserved for future use if the standard is updated

▶ Opcode

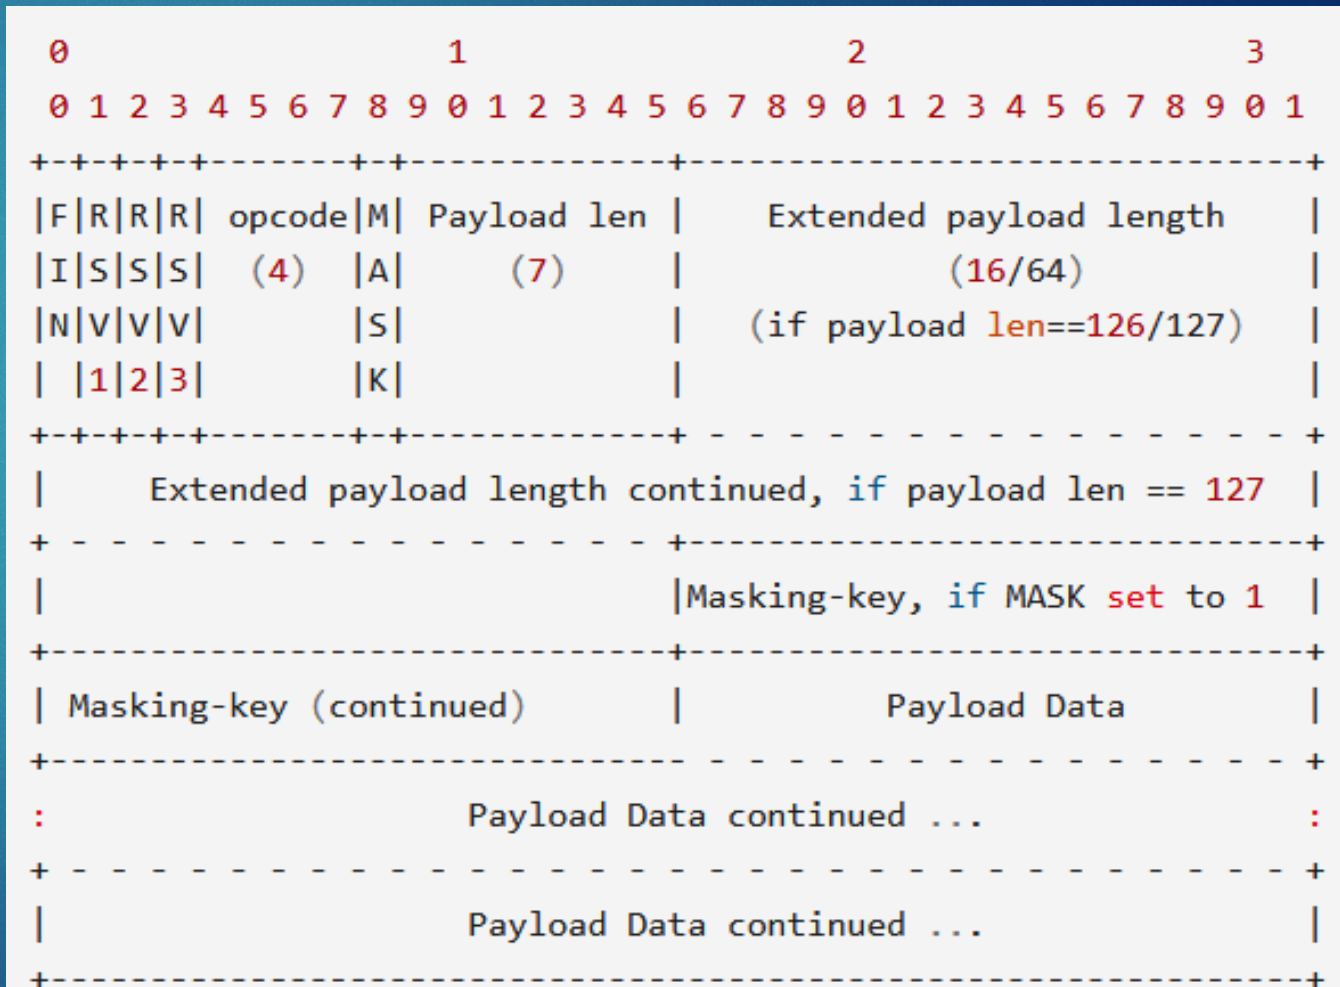
- ▶ Is it a text message, binary, etc?
- ▶ For the most part we really don't care about the opcode, but it would be good to check to make sure that it is a 0x1 (ie: a text message). [Note: An opcode of 0x2 means a binary message.]
- ▶ Note, opcode of 0x8 means the client is about to close the connection...
 - ▶ Ie, the browser window is closing



Steps to read the packet...

19

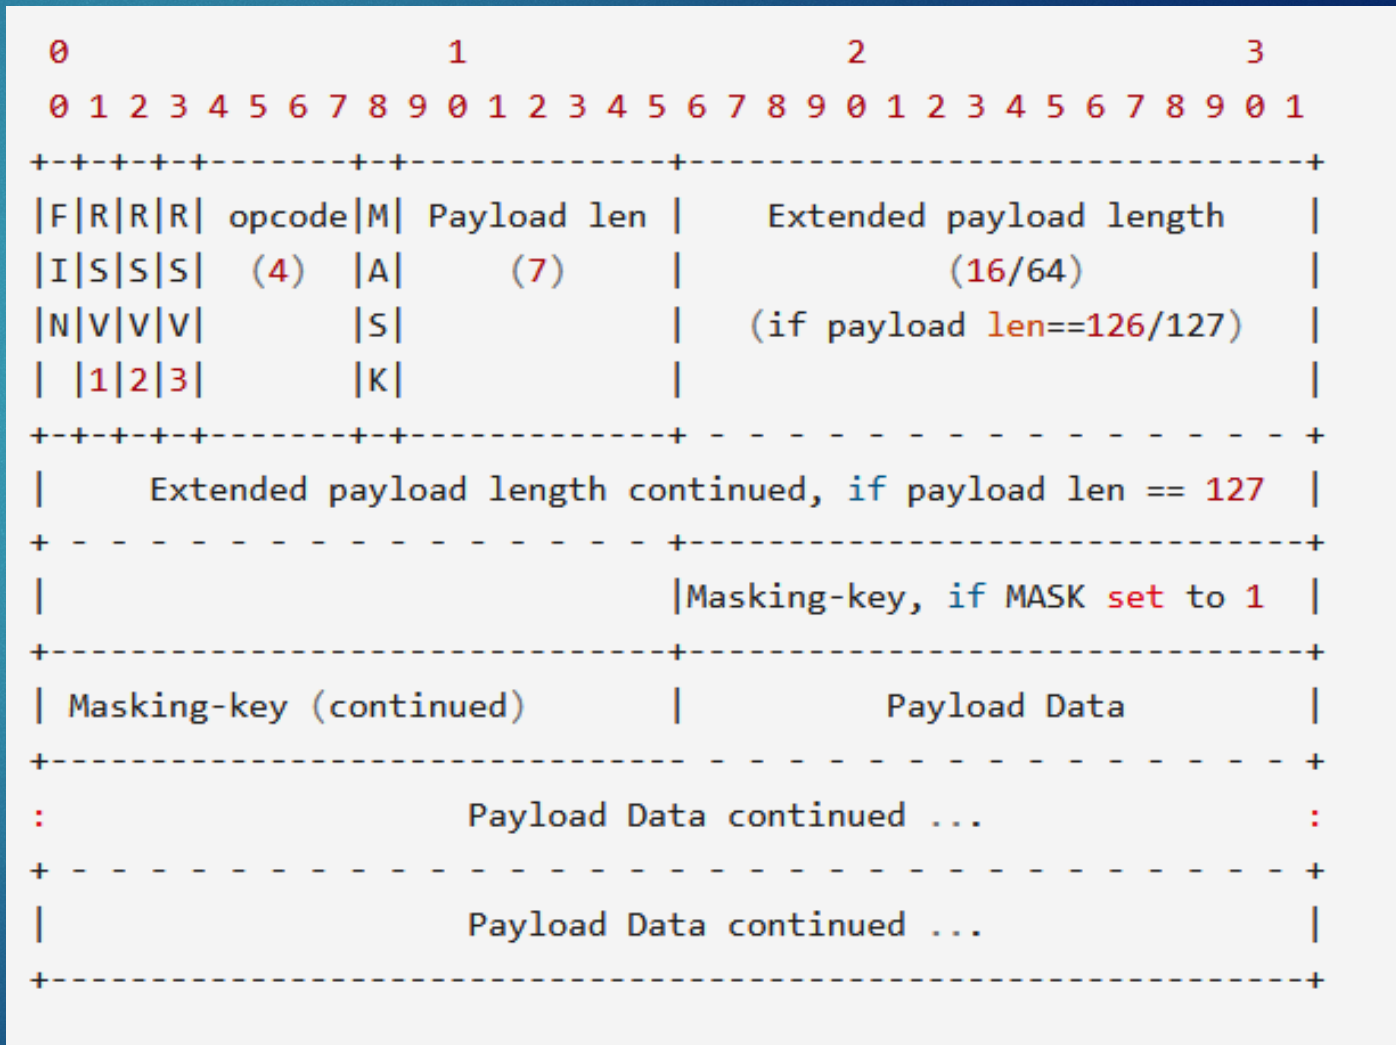
- ▶ How big is this entire message? How many bytes can we read from the socket?
 - ▶ We don't know... yet.
- ▶ How big is the header?
 - ▶ Could be several sizes...
- ▶ How much of the header do we know will always be the same size?
 - ▶ First 2 bytes.
- ▶ Once we have the first 2 bytes, how much more "header" do we read?
 - ▶ Depends on the length field.
 - ▶ Read 0, or 2, or 8 bytes more.
- ▶ Are we at the payload yet?
 - ▶ Depends on the mask bit.
 - ▶ If it is set to 1, read how much more?
 - ▶ 4 bytes.
- ▶ Now we are at the payload. 😊



In Java:

20

- ▶ Get the InputStream from the socket...
- ▶ DataInputStream wraps a normal InputStream (like a Scanner does).
- ▶ What does the Scanner do for us?
 - ▶ Makes it easy for us read strings, ints, etc.
- ▶ DataInputStream makes it easy for us to read groups of bytes (binary data).
 - ▶ readNBytes(X)
 - ▶ readFully(byte [])
 - ▶ readShort()
 - ▶ readLong()
 - ▶ readInt()



Incoming WebSocket Message

21

- ▶ We have finished the handshake and now it is time to read a real WebSocket message from the client socket.. Message looks like this...

```
<- 0x82 0xFE 0x0100 0xABCD ...
```

Remember, http packets look like this:

```
<- GET /file HTTP/1.1\r\nHost: google.com\r\n\r\n
```

- ▶ What does this mean?

- ▶ Perhaps it is easier if we look at it in binary?

```
<- 1000 0010 1111 1110 0000 0001 0000 0000 1010 1011 1100 1101 .....
```

- ▶ What do each of those bits mean? Let's look at them one byte at a time...

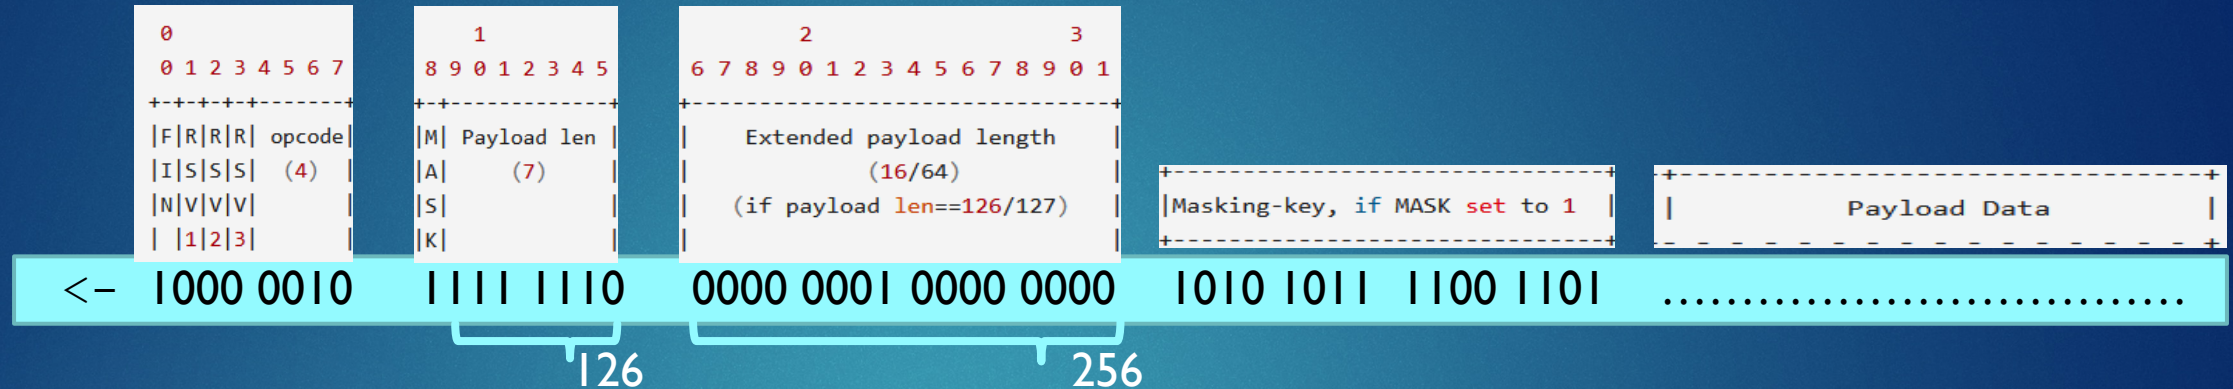
- ▶ But first, do we know how long the entire message is?

- ▶ Not yet! So how many bytes can we safely read to get started understanding / parsing this message?

▶ 2

Incoming WebSocket Message

22



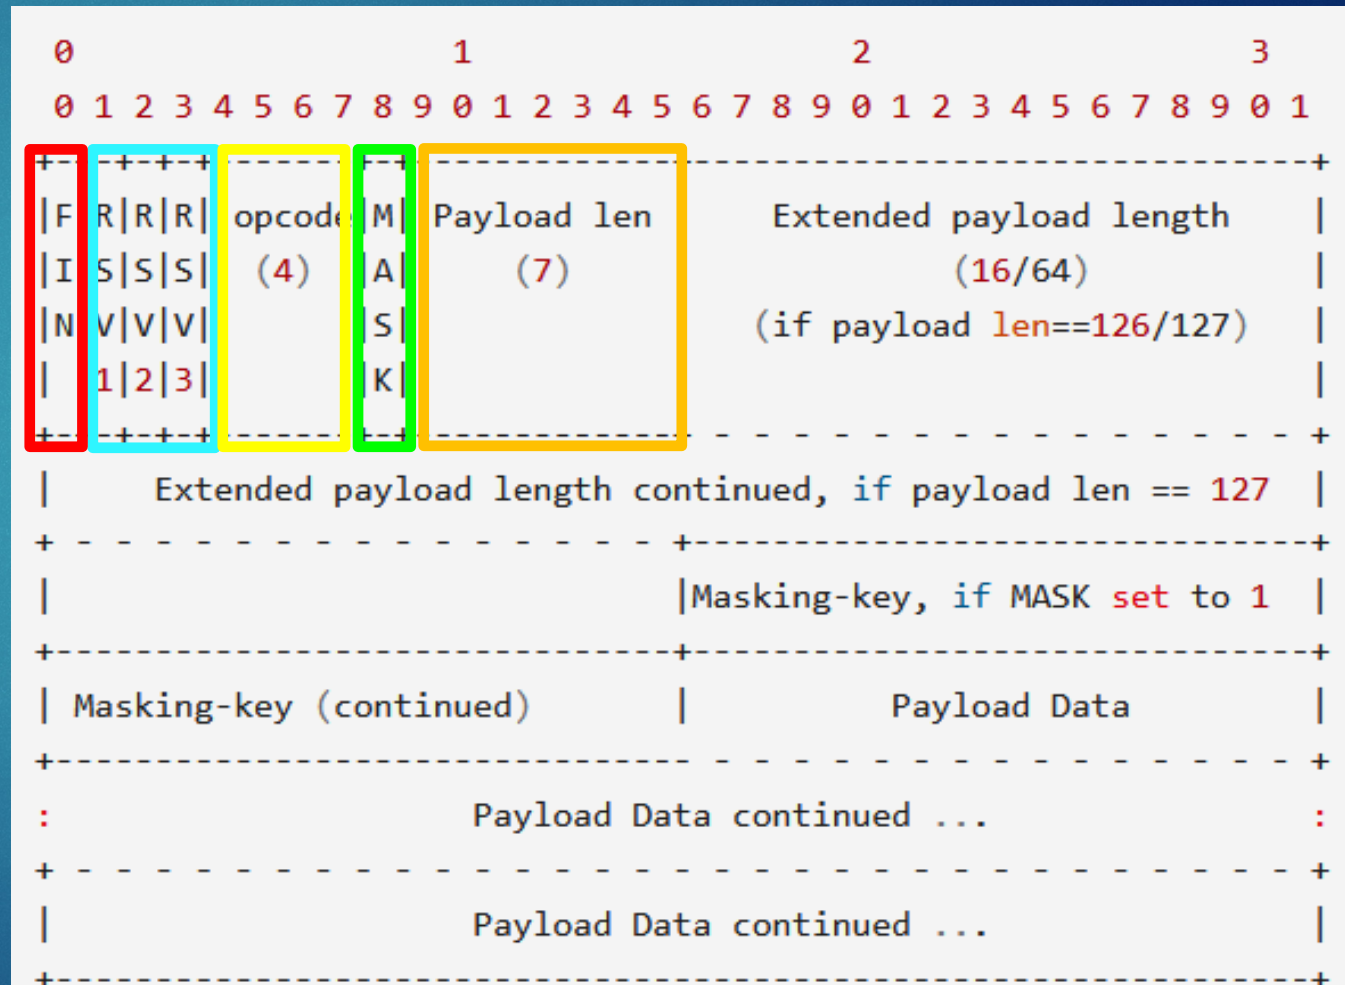
- ▶ Read 2 bytes... what does the 1st byte tell us?
 - ▶ Does looking at the spec help?
 - ▶ Fin bit is true – this is the only packet in this msg.
 - ▶ Opcode of 2 – means binary message.
- ▶ What does the 2nd byte tell us?
 - ▶ Spec?
 - ▶ Message will be masked.
 - ▶ Length is 126... which means?
 - ▶ Spec?
 - ▶ The next two bytes are the real length of the payload.
 - ▶ So the real length is?
- ▶ 256 bytes.
- ▶ What's next after the length fields?
 - ▶ Spec?
 - ▶ The mask... which is?
 - ▶ 0 x A B C D
- ▶ Finally, read the rest of the message... What is this called?
 - ▶ Payload – how many bytes?
 - ▶ 256

- | 0 | | | | | | | | | | 1 | | | | | | | | | | 2 | | | | | | | | | | 3 | | | | | | | | | |
|--------------------------------|----------------------------|---|---|---|---|---|---|---|---|--|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|---|
| 0 | 1 | 2 | 3 | 4 | 5 | 6 | 7 | 8 | 9 | 0 | 1 | 2 | 3 | 4 | 5 | 6 | 7 | 8 | 9 | 0 | 1 | 2 | 3 | 4 | 5 | 6 | 7 | 8 | 9 | 0 | 1 | 2 | 3 | 4 | 5 | 6 | 7 | 8 | 9 |
| +-----+-----+-----+-----+ | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | |
| F R R R opcode M Payload len | | | | | | | | | | Extended payload length | | | | | | | | | | | | | | | | | | | | | | | | | | | | | |
| I S S S (4) A (7) | | | | | | | | | | (16/64) | | | | | | | | | | | | | | | | | | | | | | | | | | | | | |
| N V V V S | | | | | | | | | | (if payload len==126/127) | | | | | | | | | | | | | | | | | | | | | | | | | | | | | |
| 1 2 3 K | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | |
| +-----+-----+-----+-----+ | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | |
| | | | | | | | | | | Extended payload length continued, if payload len == 127 | | | | | | | | | | | | | | | | | | | | | | | | | | | | | |
| +-----+-----+-----+-----+ | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | |
| | | | | | | | | | | Masking-key, if MASK set to 1 | | | | | | | | | | | | | | | | | | | | | | | | | | | | | |
| +-----+-----+-----+-----+ | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | |
| Masking-key (continued) | | | | | | | | | | Payload Data | | | | | | | | | | | | | | | | | | | | | | | | | | | | | |
| +-----+-----+-----+-----+ | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | |
| : | Payload Data continued ... | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | : | | |
| +-----+-----+-----+-----+ | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | |
| : | Payload Data continued ... | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | : | | |
| +-----+-----+-----+-----+ | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | |

Responding

24

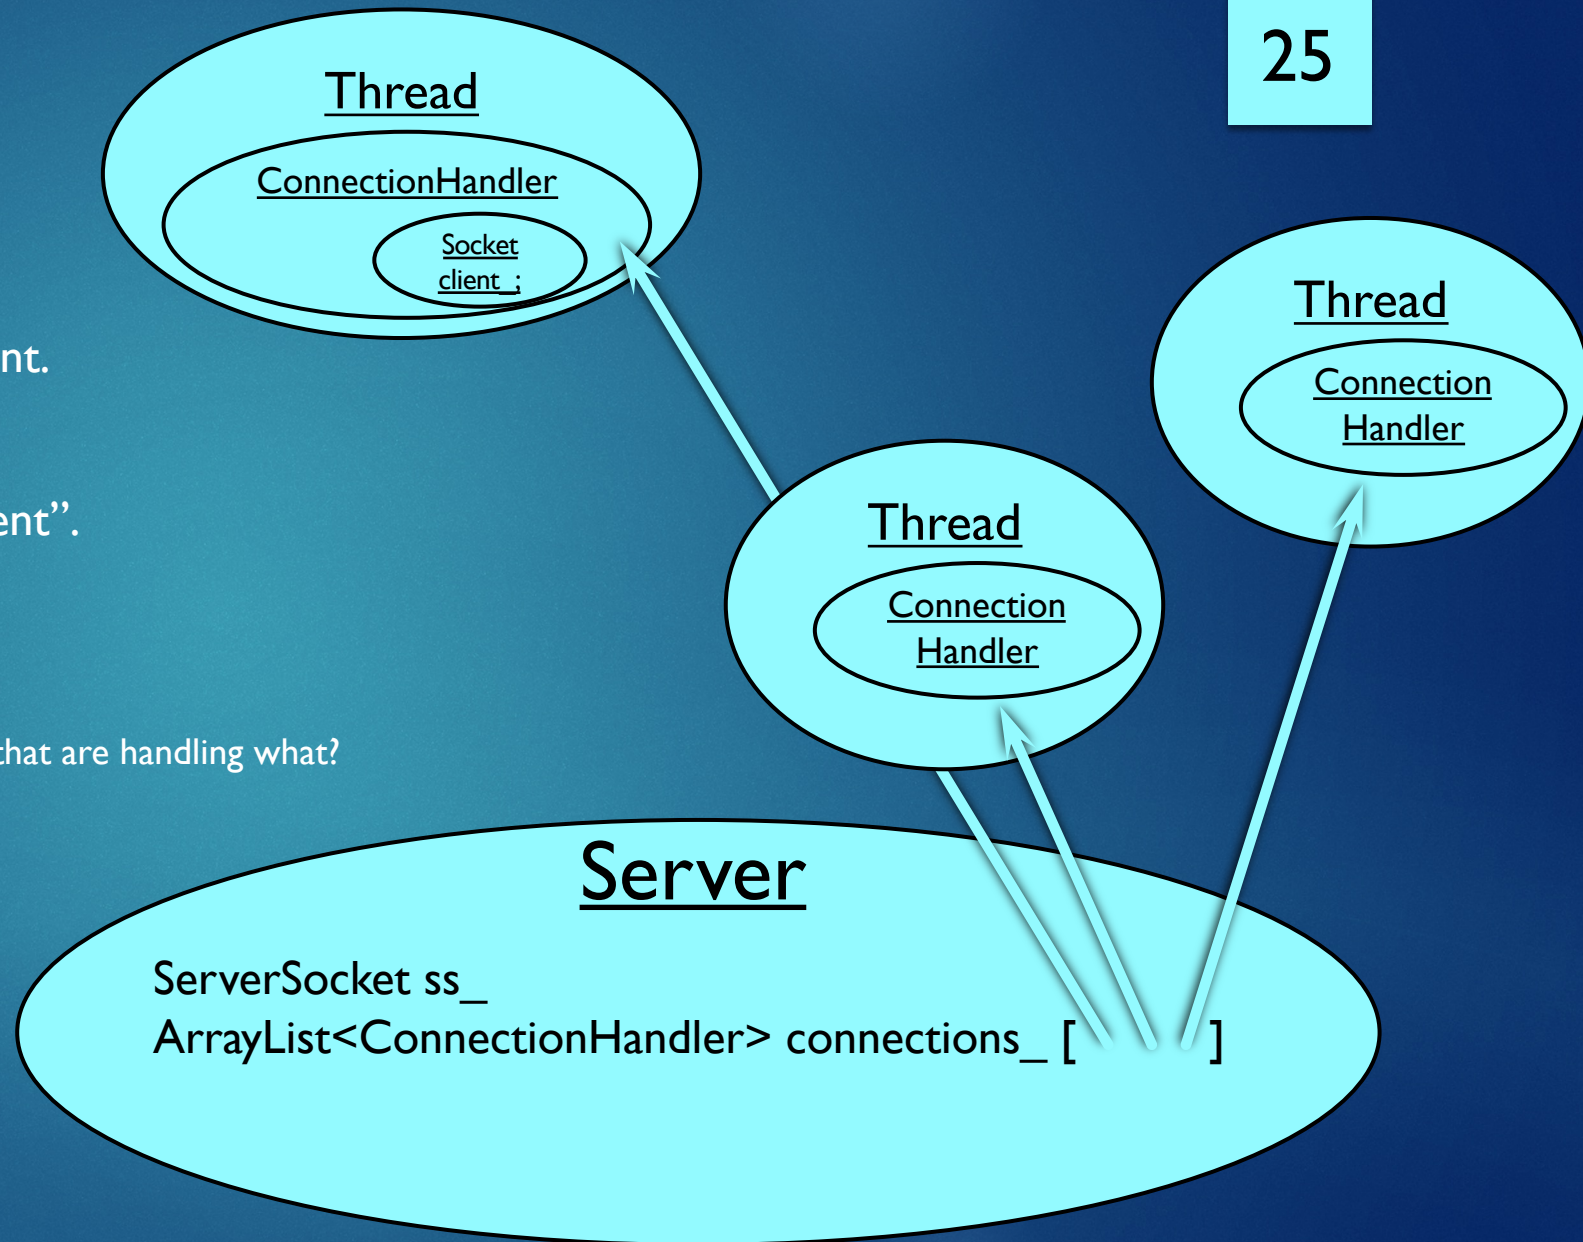
- ▶ Fill in the information need on based on the specification to the right:
- ▶ 1000 0001 0000 1100 <payload>
- ▶ What do all those 1s and 0s mean?
- ▶ 1000 0001 0000 1100 <payload>
- ▶ 1 – Final packet in this message.
- ▶ 0001 – Opcode. Value? Means?
 - ▶ 1, Text Message
- ▶ 0 – No masking... why?
 - ▶ Spec says messages from server to client are not masked.
- ▶ 010 1100 – Length. Value?
 - ▶ 122 – { "type" : "message", "user" : "Davison", "room" : "testroom", "message" : "How are you?" }
- ▶ Payload?
 - ▶ Just ASCII { "type" : "message", "user" : "Davison", "room" : "testroom", "message" : "How are you?" }



Responding

25

- ▶ Sends back a* response... ???
 - ▶ Only one?
 - ▶ A message to every connected client.
 - ▶ For now.
 - ▶ What determines a “connected client”.
 - ▶ The ConnectionHandler Thread...
 - ▶ Always?
 - ▶ No – only ConnectionHandlers that are handling what?
 - ▶ WebSocket requests
 - ▶ Who keeps track of all the ConnectionHandlers?
 - ▶ The Server – How? Draw a bubble diagram of the program.



Testing

- ▶ You can use the Spec – Section 5.7:
 - ▶ Unmasked text message:
 - ▶ 0x81 0x05 0x48 0x65 0x6c 0x6c 0x6f
 - ▶ contains a body of "Hello"
 - ▶ Masked text message:
 - ▶ 0x81 0x85 0x37 0xfa 0x21 0x3d 0x7f 0x9f 0x4d 0x51 0x58
 - ▶ contains a body of "Hello"
 - ▶ 256 bytes binary message in a single unmasked frame
 - ▶ 0x82 0x7E 0x0100 [256 bytes of binary data]
 - ▶ 64 KiB binary message in a single unmasked frame
 - ▶ 0x82 0x7F 0x0000000000010000 [65536 bytes of binary data]
- ▶ Can use your client to send a (WebSocket) message to the server and decode it.

0	1	2	3	4	5	6	7	8	9	0	1	2	3	4	5	6	7	8	9	0	1	2	3	4	5	6	7	8	9	0	1
+	+	+	+	+	+	+	+	+	+	+	+	+	+	+	+	+	+	+	+	+	+	+	+	+	+	+	+	+	+	+	+
	F		R		R		R		R		opcode		M		Payload	len		Extended payload length													
	I		S		S		S		S		(4)		A		(7)		(16/64)														
	N		V		V		V		V				S				(if payload len==126/127)														
			1		2		3						K																		
+	+	+	+	+	+	+	+	+	+	+	+	+	+	+	+	+	+	+	+	+	+	+	+	+	+	+	+	+	+	+	+

Echo Server

27

- ▶ Client sends a WebSocket message
- ▶ Server reads it, decodes it, and sends it back
- ▶ Why are we using a Threaded server?
 - ▶ Because the WebSocket stays open “forever” on the server sitting in a loop:

```
while( true ) {  
    // Read (WebSocket) message.  
    // Respond to message... with a WebSocket message.  
}
```

- ▶ Note, the server should continue to serve normal HTTP requests also.

Wednesday Assignments

- ▶ Code Review – Multithreaded Web Server
- ▶ Assignment – WebSocket Echoes

~ Fin ~